

# Stable Arm-Swing Humanoid Locomotion under Upper-Body Perturbations

A staged reward-routing curriculum preserves full-body walking, introduces coordinated shoulder-pitch arm swing, and evaluates recovery under torso, shoulder, and elbow action pulses.

Course Project Team • H1 Humanoid • 19DoF full-body control • Headless evaluation in Genesis

**Robust walking is learned without freezing upper-body motion.**

The final policy survives strong upper-body perturbations while retaining visible shoulder swing and forward command tracking.

**19**  
CONTROLLED  
DOF

**32**  
PARALLEL  
ENVS

**20s**  
EVAL  
HORIZON

**0**  
FALLS AT  
AMP2

## Dynamic Evidence: from Locked Upper Body to Perturbation Recovery

### Problem formulation

Once the upper body is unlocked, humanoid locomotion is no longer only a lower-body tracking problem. Torso, shoulder, and elbow actions perturb angular momentum, contact timing, and posture, so a naive arm-swing reward can degrade the base gait.

### Contributions

- Full-body H1 control in a 19DoF action space.
- Command-routed reward families that protect walking before arm shaping.
- Upper-body action-pulse evaluation on torso, shoulder, and elbow channels.
- Qualitative temporal evidence paired with finite-horizon metrics.

### Task-to-benchmark mapping


|                    |                                                         |
|--------------------|---------------------------------------------------------|
| Forward locomotion | 0.6 m/s command tracking under stable gait.             |
| Upper body motion  | Visible shoulder-pitch swing synchronized with walking. |
| Robustness         | Recovery after held upper-body action offsets.          |

### Design principle

Locomotion is protected before expressiveness is added. Command-derived masks activate standing, walking, and straight-walk reward terms only in compatible regimes.

### A. Locked upper-body baseline


stable gait, upper body constrained



The policy first preserves stable forward walking with the upper body held near default posture.

### B. Coordinated arm swing


visible shoulder-pitch swing



After arm-swing shaping, shoulder motion becomes visible while the walking rhythm is maintained.

### C. Upper-body perturbation recovery

pulse disturbance and recovery



Upper-body action offsets disturb torso and arms; the policy continues walking after the pulse window.

FIGURE 1 • TEMPORAL ROLLOUTS

### Evaluation protocol

We evaluate 32 parallel H1 humanoids for a 20 s horizon. Upper-body action pulses are applied to torso, shoulder, and elbow channels while the forward command remains fixed at 0.6 m/s.

### Metrics reported

Survival counts environments without non-timeout resets. Tracking error measures forward velocity deviation. Max tilt uses the projected-gravity tilt proxy and exposes perturbation-induced posture excursions.

### Result interpretation

Recovery success is not used as the headline baseline comparison because a policy may briefly satisfy the recovery criterion after some pulses while still repeatedly resetting. Survival, falls, and max tilt are more decisive.

### Reading the rows

Row A verifies the stable walking foundation. Row B shows visible gait-synchronized arm swing. Row C shows a perturbed upper body returning to forward walking.


Temporal strips compare stable locked-upper-body walking, gait-synchronized arm swing, and recovery after an injected upper-body action pulse. The qualitative claim is temporal: the robot continues walking while upper-body behavior changes from constrained, to coordinated, to externally perturbed.

## Method: Staged Reward Routing

FIGURE 2 • METHOD

### Staged Reward Routing for Full-Body Arm-Swing Locomotion

Reward terms are activated by command regimes so arm-swing objectives do not destabilize the base gait.



#### Routed objective


$$R_t = R_{\text{shared}} + m_{\text{stand}} R_{\text{stand}} + m_{\text{walk}} R_{\text{walk}} + m_{\text{straight}} R_{\text{arm}} - C_{\text{upper}}$$

- shared rewards stay active: velocity, heading, contact, smoothness
- arm-swing terms activate mainly during straight walking

The staged protocol preserves the locomotion solution before strengthening upper-body expressiveness. Robustness is reported separately in the perturbation results, not mixed into the reward-routing pipeline.


The controller is trained in the full 19DoF H1 action space. Stable locomotion is protected first; command masks then activate stand, walk, and straight-walk reward families so arm-swing objectives appear mainly when they are compatible with forward walking.

## Quantitative Robustness under Upper-Body Pulses



| POLICY / SETTING   | AMP | SURVIVAL | MEAN TIME | FALLS | TRACK ERR. | MAX TILT |
|--------------------|-----|----------|-----------|-------|------------|----------|
| Robust policy      | 0.0 | 100%     | 20.00s    | 0     | 0.0228     | 4.67°    |
| Robust policy      | 1.0 | 100%     | 20.00s    | 0     | 0.0225     | 7.81°    |
| Robust policy      | 2.0 | 100%     | 20.00s    | 0     | 0.0245     | 9.12°    |
| Arm-swing baseline | 2.0 | 0%       | 5.48s     | 95    | 0.1401     | 89.93°   |

#### Robust policy under increasing perturbation amplitude



The amp0 run uses the same perturbation scheduling with zero magnitude. At amp2, the robust policy retains arm motion: shoulder-pitch swing amplitude 0.2536 rad and elbow sagittal separation 0.1165 m, so robustness is not obtained by simply freezing the upper body.

## Experimental Story

- Locomotion foundation**  
Full-body H1 walks forward while the upper body is constrained to preserve a reliable gait.
- Command-routed preparation**  
Stand, walk, and straight-walk masks reduce interference between posture, gait, and arm objectives.
- Coordinated arm swing**  
Shoulder-pitch phase and endpoint-sagittal terms make the upper body visibly participate in walking.
- Perturbation-robust policy**  
Held action offsets on torso, shoulder, and elbow channels test whether the gait recovers.

Takeaway: unlocking upper-body joints is not enough. Stable arm-swing locomotion emerges when walking is protected first, arm coordination is routed by command regime, and robustness is evaluated with explicit upper-body disturbances.

### Limitations

Evidence is simulation-only and short-horizon. Upper-body naturalness is shaped by hand-designed phase and amplitude rewards rather than human motion priors. Future work should add angular-momentum analysis, energy cost, longer horizons, and cleaner camera evaluation.

### Conclusion

The final controller demonstrates stable forward walking, visible arm swing, and bounded recovery behavior under explicit upper-body disturbances, forming a compact paper-style story from curriculum design to quantitative robustness.